

Programming Assignment 5: Machine Translation

CSE 256: Statistical NLP: Spring 2022

University of California, San Diego

Released: May 09, 2022

Due: May 23, 2022

For this assignment you will develop models for machine translation. The goal is to implement two translation models, IBM model 1 and IBM model 2, and apply these models to predict English/Spanish word alignments. Both models estimate the conditional probability of a foreign sentence $f_1 \dots f_m$ and an alignment $a_1 \dots a_m$ given a particular English sentence $e_1 \dots e_l$ and length m . The simpler model, IBM model 1, defines the conditional probability with a single set of parameters t as

$$P(f_1 \dots f_m, a_1 \dots a_m | e_1 \dots e_l, m) = \frac{1}{(l+1)^m} \prod_{i=1}^m t(f_i | e_{a_i})$$

The richer model, IBM model 2, defines the same conditional probability with two sets of parameters t and q

$$P(f_1 \dots f_m, a_1 \dots a_m | e_1 \dots e_l, m) = \prod_{i=1}^m q(a_i | i, l, m) t(f_i | e_{a_i})$$

1 Preliminaries

1.1 Training Data

The main task of this assignment is to estimate these parameters from data. We estimate the parameters from a parallel corpus of aligned sentences. Each entry in the corpus contains two versions of the same sentence, one in English and one in Spanish. The parallel corpus we are using is a subset of the Europarl corpus [Koehn, 2005] accessed from <http://www.statmt.org/europarl/> which contains aligned sentences from the proceedings of the European Parliament.

The training corpus is split into two files, `corpus.en` and `corpus.es`, which contain English and Spanish sentences respectively. The i -th sentence in the English file is a translation of the i -th sentence in the Spanish file. The files are in UTF-8, contain one sentence per line, and words are separated by a space.

```
> head -n 3 corpus.en
resumption of the session
i declare resumed the session of the european parliament adjourned on ...
Although , as you will have seen , the dreaded ' millennium bug ' failed ...
```

```
> head -n 3 corpus.es
reanudacion del periodo de sesiones
declaro reanudado el periodo de sesiones del parlamento europeo ...
como todos han podido comprobar , el gran " efecto del año 2000 " no se ...
```

1.2 IBM Model 1

The full algorithm for IBM model 1 is given on page 21 in the Collins notes <http://www.cs.columbia.edu/~mcollins/ibm12.pdf>. We recommend closely reading those notes for a full

description of the problem. The core portion of the algorithm is summarized here.

Recall that IBM model 1 only has word translation parameters $t(f|e)$ which can be interpreted as the conditional probability of generating a foreign word f from an English word e (or from NULL).

We can estimate $t(f|e)$ using the EM algorithm, which iterates through the parallel corpus repeatedly. For the k -th sentence pair and each index i in the foreign sentence $f^{(k)}$ and j in the English sentence $e^{(k)}$, we define

$$\delta(k, i, j) = \frac{t(f_i^{(k)}|e_j^{(k)})}{\sum_{j=0}^{l_k} t(f_i^{(k)}|e_j^{(k)})}$$

where l_k is the length of the English sentence. The delta function is used to update the *expected counts* for the current iteration:

$$\begin{aligned} c(e_j^{(k)}, f_i^{(k)}) &\leftarrow c(e_j^{(k)}, f_i^{(k)}) + \delta(k, i, j) \\ c(e_j^{(k)}) &\leftarrow c(e_j^{(k)}) + \delta(k, i, j) \end{aligned}$$

After each iteration through the parallel corpus, we revise our estimate for t parameters:

$$t(f|e) = \frac{c(e, f)}{c(e)}$$

for all possible foreign words f and English words e (and NULL).

1.3 IBM Model 2

The complete algorithm for estimating IBM model 2 parameters is given on page 13 of the Collins notes <http://www.cs.columbia.edu/~mccollins/ibm12.pdf>. Be sure to understand the details of this algorithm before beginning to code the assignment.

Note that IBM model 2 extends the implementation of the EM algorithm for IBM model 1. The main additional step is adapting the delta function to include $q(j|i, l, m)$

$$\delta(k, i, j) = \frac{q(j|i, l_k, m_k) t(f_i^{(k)}|e_j^{(k)})}{\sum_{j=0}^{l_k} q(j|i, l_k, m_k) t(f_i^{(k)}|e_j^{(k)})}$$

and computing expected counts $c(j|i, l, m)$ and $c(i, l, m)$

After each iteration through the corpus we re-estimate the $t(f_i, e_j)$ parameters as before and add new updates for the q parameters:

$$q(j|i, l, m) = \frac{c(j|i, l, m)}{c(i, l, m)}$$

1.4 Evaluation

We provide aligned development and test data, `dev.en` / `dev.es` in order to evaluate alignment accuracy. The format of these files is identical to the corpus training files, one sentence per line with aligned sentences.

We also provide a file, `dev.key`, which contains the manually annotated gold alignments for the development sentences. The file contains lines of the form

```
...
6 7 3
6 7 4
6 8 5
...
```

Each line specifies a word alignment of the form `SentenceIndex EnglishIndex ForeignIndex`. The first line of the example specifies that in the 6th pair of sentences the 7th English word is aligned to the 3rd Spanish word. Note that the gold alignments do not include NULL word alignments, all three indices start from 1, and there are no blank lines between sentences.

As an example, suppose the 6th sentence pair is:

- e = Do we want to make it happen ?
- f = ¿ Deseamos hacerla realidad ?

The gold word alignment, 6 7 3, indicates that the English word “happen” is aligned with the Spanish word “hacerla”. The full alignment of the sentence pair is

	¿	Deseamos	hacerla	realidad	?
Do					
we					
want					
to					
make					
it					
happen					
?					

For all of the questions in this assignment you should output predicted alignments in this format. We have included a `script_eval_alignment.py` to evaluate the accuracy of these alignments. Assuming that you have written alignments to a file `dev.out`, running `python eval_alignment.py dev.key dev.out` will give the F-Score of the predicted alignments compared to the gold alignments.

Note: The alignments given in the key file often map single foreign words to multiple English words, whereas our models do not. For this reason it is impossible to get perfect recall or F-Score on this dataset. For parts 1 and 2 we ignore this issue.

2 Part 1 - IBM Model 1 (6 Points)

The first problem is to estimate the parameters of IBM model 1 using `corpus.en` and `corpus.es` as input. Before implementing please read the following notes carefully

- Your implementation should only store t parameters for possible pairs of foreign and English words, i.e. words that occur together in some parallel translation, and the special English word NULL. We recommend implementing t as a collection of sparse maps for each English word e where the keys of the map are the words f that occur with e and the values are the corresponding parameters $t(f|e)$.

Using a non-sparse implementation will likely result in memory issues.

- In the initialization step, set $t(f|e)$ to be the uniform distribution over all foreign words that could be aligned to e in the corpus. More specifically

$$t(f|e) = \frac{1}{n(e)}$$

where $n(e)$ is the number of different words that occur in any translation of a sentence containing e . Note that the special English word NULL can be aligned to any foreign word in the corpus.

- Starting from the initial $t(f|e)$ parameters, run 5 iterations of the EM algorithm for IBM model 1 (this may take a while). When your model is completed, save your t parameters to a file, as you will need them for the next problem.
- Finally, use your model to find alignments for the development sentence pairs `dev.en / dev.es`. For each sentence, align each foreign word f_i to the English word with the highest $t(f|e)$ score,

$$a_i = \arg \max_{j \in 0 \dots l} t(f_i|e_j)$$

Write your alignments to a file in the format of `dev.key` described above. Check the accuracy with `eval_alignments.py`.

- Your report should include a brief summary of your IBM Model 1 implementation and any specific challenges or issues that came up. If you could not get your implementation working, describe where you got stuck.
- The expected F1-Score is 0.42.

3 Part 2 - IBM Model 2 (6.5 Points)

We now extend the alignment model to IBM model 2 by adding alignment parameters $q(j|i, l, m)$

- Initialize the q parameters to the uniform distribution over all j for each i, l , and m .

$$q(j|i, l, m) = \frac{1}{l+1}$$

You only need to store parameters for pairs of sentence lengths l and m that occur in the corpus.

- To initialize the $t(f|e)$ parameters, use the last set of parameters (after 5 iterations) produced by your implementation of IBM model 1.
- Run 5 iterations of EM for IBM model 2.
- As before, use the model to compute alignments for the development sentence pairs in the corpus. For each foreign word f_i , the best alignment is

$$a_i = \arg \max_{j \in 0 \dots l} q(j|i, l, m) t(f_i|e_j)$$

Write your alignments to a file in the format of `dev.key` described above. Check the accuracy with `eval_alignments.py`.

- Your report should include a brief summary of your IBM Model 2 implementation and any specific challenges or issues that came up. If you could not get your implementation working, describe where you got stuck.
- The expected F1-Score is 0.45.

4 Report Guideline

4.1 IBM Model 1

- Description of IBM Model 1 - What are IBM Models used for? What are the limitations?
- Description of EM Algorithm
 - What are the strengths and weaknesses?
- Method Overview
 - Provide a high-level description of your implementation
- Results
 - Report your F1 score. It should match the expected result. If not, partial credits will be given if you discuss some challenges and issues in your implementation.
- Discussion - Show how F1 score changes after each iteration; comment on the findings.

4.2 IBM Model 2

- Description of IBM Model 2
 - Why is it better than IBM Model 1? What are the limitations?
- Method Overview
 - Provide a brief, high-level description of your implementation
- Results
 - Report your F1 score. It should match the expected result. If not, partial credits will be given if you discuss some challenges and issues in your implementation.
- Discussion
 - Show how F1 score changes after each iteration; comment on the findings
 - Use the type of table/graph as in the slides or in the instructions, show a correctly aligned example + a misaligned example. Discuss the examples.

5 Submission Instructions

Submit your work on Gradescope.

- **Code:** You will submit your code together with a neatly written README file to instruct how to run your code with different settings. We assume that you always follow good practice of coding (commenting, structuring), and these factors are not central to your grade.
- **Report:** Submit your report, it should be **four pages long, or less, in pdf** (reasonable font sizes).

6 Acknowledgments

Adapted with minor changes from a similar assignment by Michael Collins.

7 References

Philipp Koehn. Europarl: A parallel corpus for statistical machine translation. In MT summit, volume 5, 2005.

Patrik Lambert, Adria De Gispert, Rafael Banchs, and Jose B Marino. Guidelines for word alignment evaluation and manual alignment. Language Resources and Evaluation, 39(4):267–285, 2005.